

# DATA MODELS AND DATABASES

BILL HOWE, PHD

DIRECTOR OF RESEARCH, SCALABLE DATA ANALYTICS

UNIVERSITY OF WASHINGTON ESCIENCE INSTITUTE




# HOW DO WE STORE DATA?



# HOW DO WE STORE DATA?

Depth (m)	Conc NO2 (nM)	Conc NH4 (nM)
5	0	35.67
35	125	181.89
40	110	
45	165	
50	125	
60	290	
70	445	0
85	0	
105	0	
300	0	0
GoFlo 0055	70	455 16.06
GoFlo 0052	70	445 3.51



## ANNOTATIONSUMMARY-COMBINEDORFANNOTATION16 Phaeo\_genome

What is the *data model*?

###query	length	COG hit #1	e-value #1	identity #1	score #1	hit length #1	description #1
chr_4[480001-580000].287	4500						
chr_4[560001-660000].1	3556						
chr_9[400001-500000].503	4211	COG4547	2.00E-04	19	44.6	620	Cobalamin biosynthesis protein
chr_9[320001-420000].548	2833	COG5406	2.00E-04	38	43.9	1001	Nucleosome binding factor SPN
chr_27[320001-404298].20	3991	COG4547	5.00E-05	18	46.2	620	Cobalamin biosynthesis protein
chr_26[320001-420000].378	3963	COG5099	5.00E-05	17	46.2	777	RNA-binding protein of the Puf
chr_26[400001-441226].196	2949	COG5099	2.00E-04	17	43.9	777	RNA-binding protein of the Puf
chr_24[160001-260000].65	3542						
chr_5[720001-820000].339	3141	COG5099	4.00E-09	20	59.3	777	RNA-binding protein of the Puf
chr_9[160001-260000].243	3002	COG5077	1.00E-25	26	114	1089	Ubiquitin carboxyl-terminal hyc
chr_12[720001-820000].86	2895	COG5032	2.00E-09	30	60.5	2105	Phosphatidylinositol kinase and
chr_12[800001-900000].109	1463	COG5032	1.00E-09	30	60.1	2105	Phosphatidylinositol kinase and
chr_11[1-100000].70	2886						
chr_11[80001-180000].100	1523						

# WHAT IS A DATA MODEL?

## **Three components:**

1. Structures
2. Constraints
3. Operations

# EXAMPLES

## 1. Structures

- rows and columns?
- nodes and edges?
- key-value pairs?
- a sequence of bytes?

## 2. Constraints

- all rows must have the same number of columns
- all values in one column must have the same type
- a child cannot have two parents

## 3. Operations

- find the value of key x
- find the rows where column “lastname” is “Jordan”
- get the next N bytes


# WHAT IS A DATABASE?

*A collection of information  
organized to afford efficient  
retrieval*

# ANOTHER VIEW

“When people use the word database, fundamentally what they are saying is that the data should be self-describing and it should have a schema. That’s really all the word database means.”

-- Jim Gray, “The Fourth Paradigm”



# WHY WOULD I WANT A DATABASE?

## 1. Sharing

Support concurrent access by multiple readers and writers

## 2. Data Model Enforcement

Make sure all applications see clean, organized data

## 3. Scale

Work with datasets too large to fit in memory

## 4. Flexibility

Use the data in new, unanticipated ways



# QUESTIONS TO CONSIDER

**How is the data physically organized on disk?**

**What kinds of queries are efficiently supported by this organization, and what kinds are not?**

**How hard is it update the data, or add new data?**

**What happens when I encounter new queries that I didn't anticipate? Do I reorganize the data? How hard is that?**

# MOTIVATING RELATIONAL DATABASES

BILL HOWE, PHD

DIRECTOR OF RESEARCH, SCALABLE DATA ANALYTICS

UNIVERSITY OF WASHINGTON ESCIENCE INSTITUTE

# QUESTIONS TO CONSIDER

**How is the data physically organized on disk?**

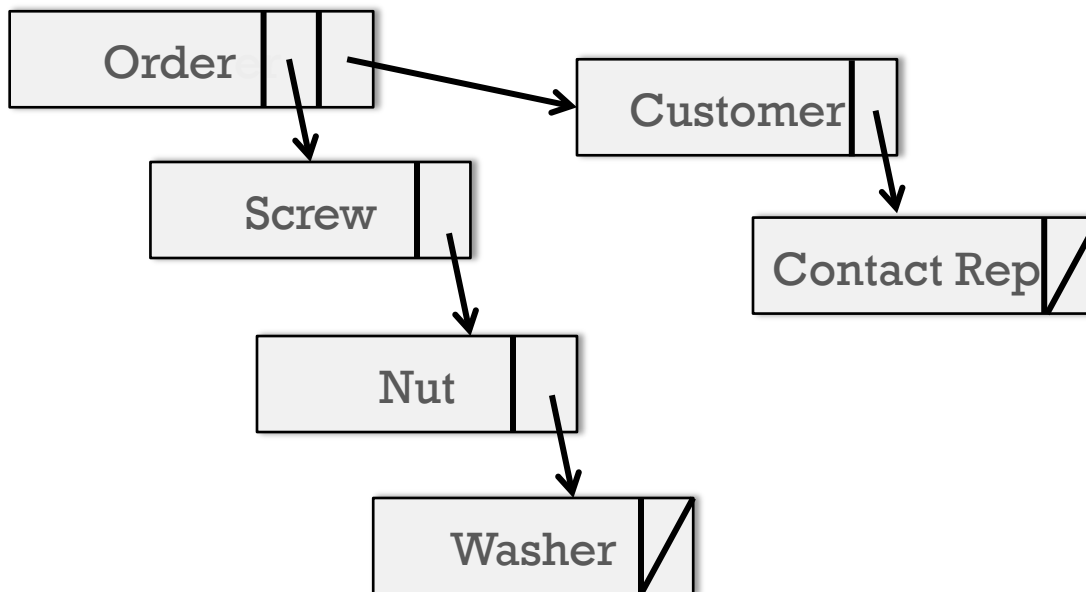
**What kinds of queries are efficiently supported by this organization, and what kinds are not?**

**How hard is it update the data, or add new data?**

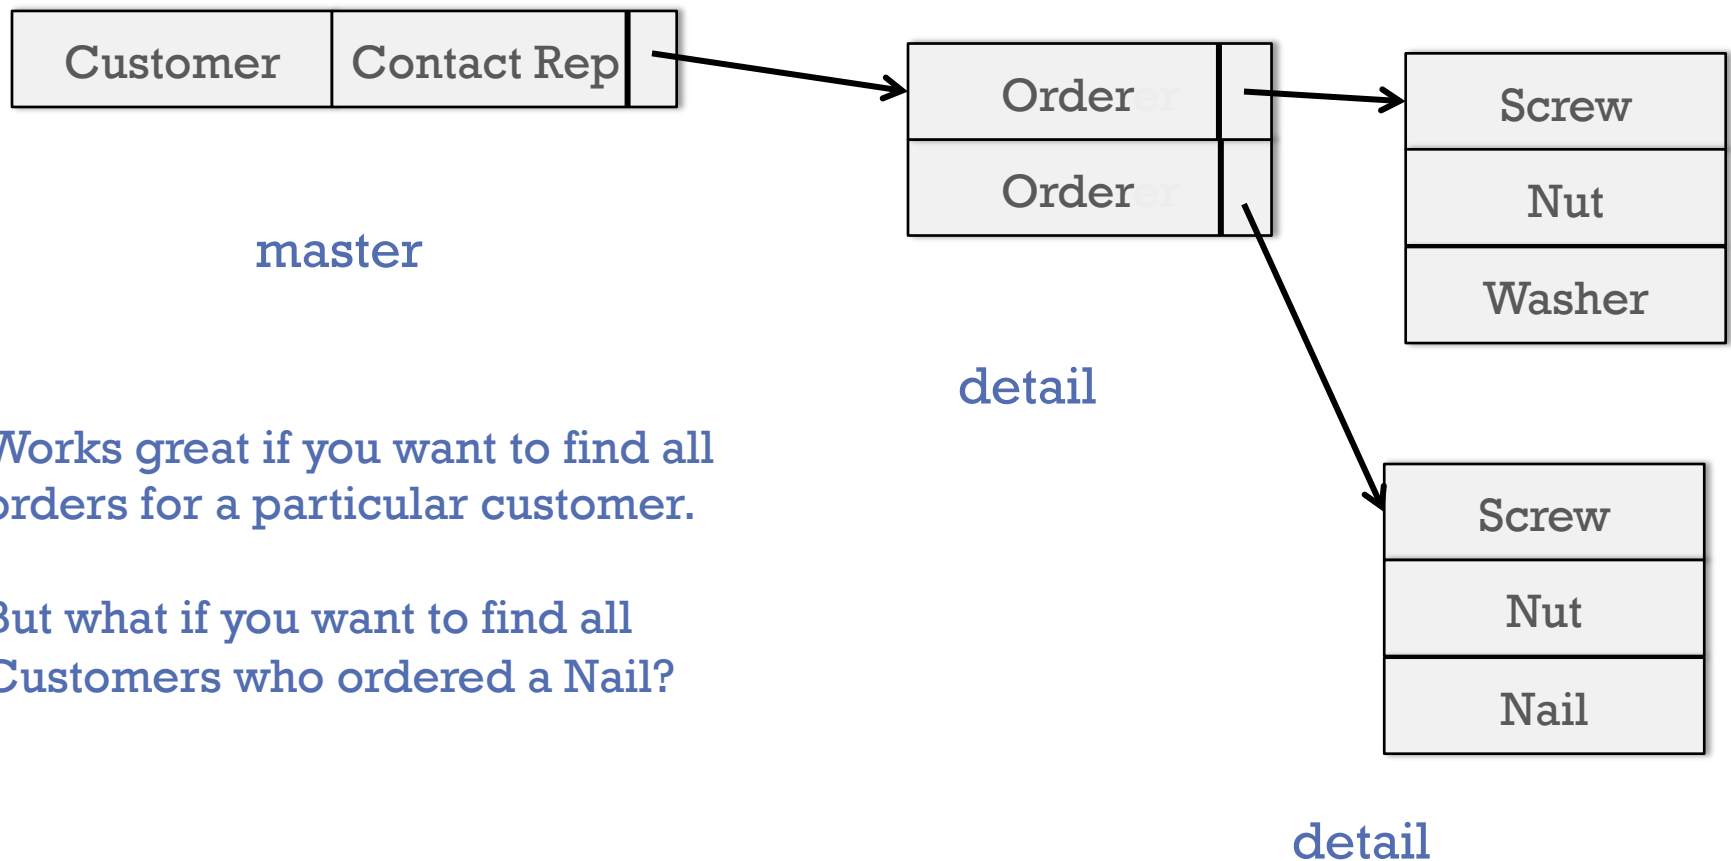
**What happens when I encounter new queries that I didn't anticipate? Do I reorganize the data? How hard is that?**

# HISTORICAL EXAMPLE: NETWORK DATABASES

*Database: A collection of information organized to afford efficient retrieval*



# HISTORICAL EXAMPLE: HIERARCHICAL DATABASES



# ONE VIEW

“Relational Database Management Systems were invented to let you use one set of data in multiple ways, including ways that are unforeseen at the time the database is built and the 1st applications are written.”

- Curt Monash, analyst/blogger

# *RELATIONAL* DATABASES (CODD 1970)

**Everything is a table**

**Every row in a table has the same columns**

**Relationships are implicit: no pointers**

Course	Student Id
CSE 344	223...
CSE 344	244...
CSE 514	255..
CSE 514	244...

Student Id	Student Name
223...	Jane
244...	Joe
255..	Susan

# ***DATABASE PHILOSOPHY***

**God made the integers;  
all else is the work of  
man.**

- Leopold Kronecker, 19<sup>th</sup> Century  
Mathematician

**Codd made relations;  
all else is the work of man.**

- Raghu Ramakrishnan, DB text book author



# RELATIONAL DATABASE HISTORY

**Pre-Relational:** if your data changed, your application broke.


Early RDBMS were buggy and slow (and often reviled), but required only 5% of the application code.

*“Activities of users at terminals and most application programs should remain unaffected when the internal representation of data is changed and even when some aspects of the external representation are changed.”*


*- Codd 1979*

**Key Ideas:** Programs that manipulate tabular data exhibit an algebraic structure allowing reasoning and manipulation independently of physical data representation

# KEY IDEA: “PHYSICAL DATA INDEPENDENCE”



*physical data independence*



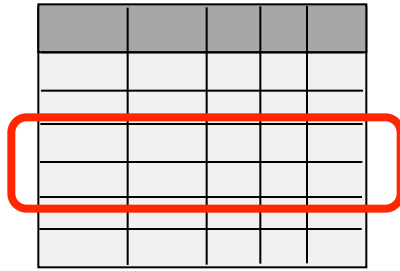
files and  
pointers

```
SELECT seq
FROM ncbi_sequences
WHERE seq =
'GATTACGATATTA' ;
```

---

```
f = fopen( 'table_file' );
fseek(10030440);
while (True) {
    fread(&buf, 1, 8192, f);
    if (buf == GATTACGATATTA) {
        . . .
```

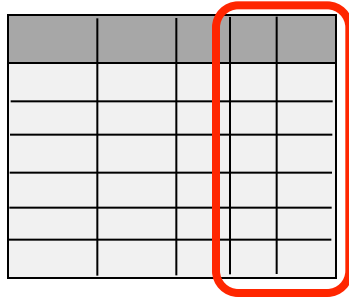
# KEY IDEA: AN *ALGEBRA OF TABLES*







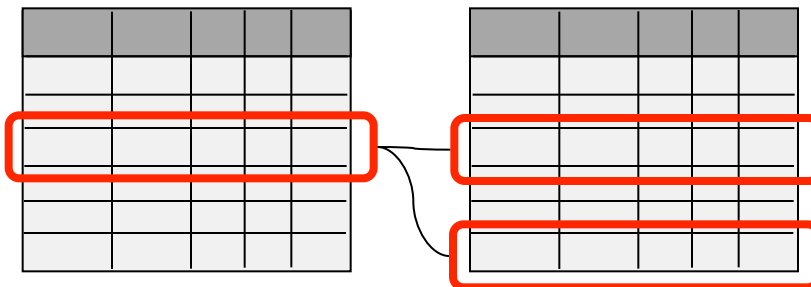
select







project











join

*Other operators: aggregate, union, difference, cross product*

# KEY IDEA: ALGEBRAIC OPTIMIZATION

$$N = ((z*2)+((z*3)+0))/1$$

## Algebraic Laws:

- 1. (+) identity:**  $x+0 = x$
- 2. (/) identity:**  $x/1 = x$
- 3. (\*) distributes:**  $(n*x+n*y) = n*(x+y)$
- 4. (\*) commutes:**  $x*y = y*x$

Apply rules **1, 3, 4, 2**:

$$N = (2+3)*z$$

**two operations instead of five, no division operator**

*Same idea works with the Relational Algebra!*

# EQUIVALENT LOGICAL EXPRESSIONS; DIFFERENT COSTS

$$\sigma_{p=\text{knows}}(R) \bowtie_{o=s} \left( \sigma_{p=\text{holdsAccount}}(R) \bowtie_{o=s} \sigma_{p=\text{accountHomepage}}(R) \right)$$

*right associative*

$$\left( \sigma_{p=\text{knows}}(R) \bowtie_{o=s} \sigma_{p=\text{holdsAccount}}(R) \right) \bowtie_{o=s} \sigma_{p=\text{accountHomepage}}(R)$$

*left associative*

$$\sigma_{p1=\text{knows} \ \& \ p2=\text{holdsAccount} \ \& \ p3=\text{accountHomepage}}(R \times R \times R)$$

*cross product*

# SAME LOGICAL EXPRESSION, DIFFERENT PHYSICAL ALGORITHMS

A = select(p=knows)  
B = select(p=holdsAccount)  
C = select(p=accountWebpage)

$hA = \text{hash}(A)$   
 $AB = \text{probe } hA \text{ with } B$

$hC = \text{hash}(C)$   
 $ABC = \text{probe } hC \text{ with } AB$


A = select(p=knows)  
B = select(p=holdsAccount)  
C = select(p=accountWebpage)

$hB = \text{hash}(B)$   
 $AB = \text{probe } hB \text{ with } A$

$hC = \text{hash}(C)$   
 $ABC = \text{probe } hC \text{ with } AB$

*Which is faster?*

# ALGEBRAIC OPTIMIZATION MATTERS




*BTC 2010 Dataset*

*3B quads*

*623 GB processed*

# PICASSO QUERY PLAN DIAGRAMS






# KEY IDEA: DECLARATIVE LANGUAGES

*Find all orders from today, along with the items ordered*

```
SELECT *  
  FROM Order o, Item i  
 WHERE o.item = i.item  
    AND o.date = today()
```



# SQL IS THE “WHAT” NOT THE “HOW”

Product(pid, name, price)

Purchase(pid, cid, store)

Customer(cid, name, city)

```
SELECT DISTINCT x.name, z.name
FROM Product x, Purchase y, Customer z
WHERE x.pid = y.pid and y.cid = y.cid and
      x.price > 100 and z.city = 'Seattle'
```

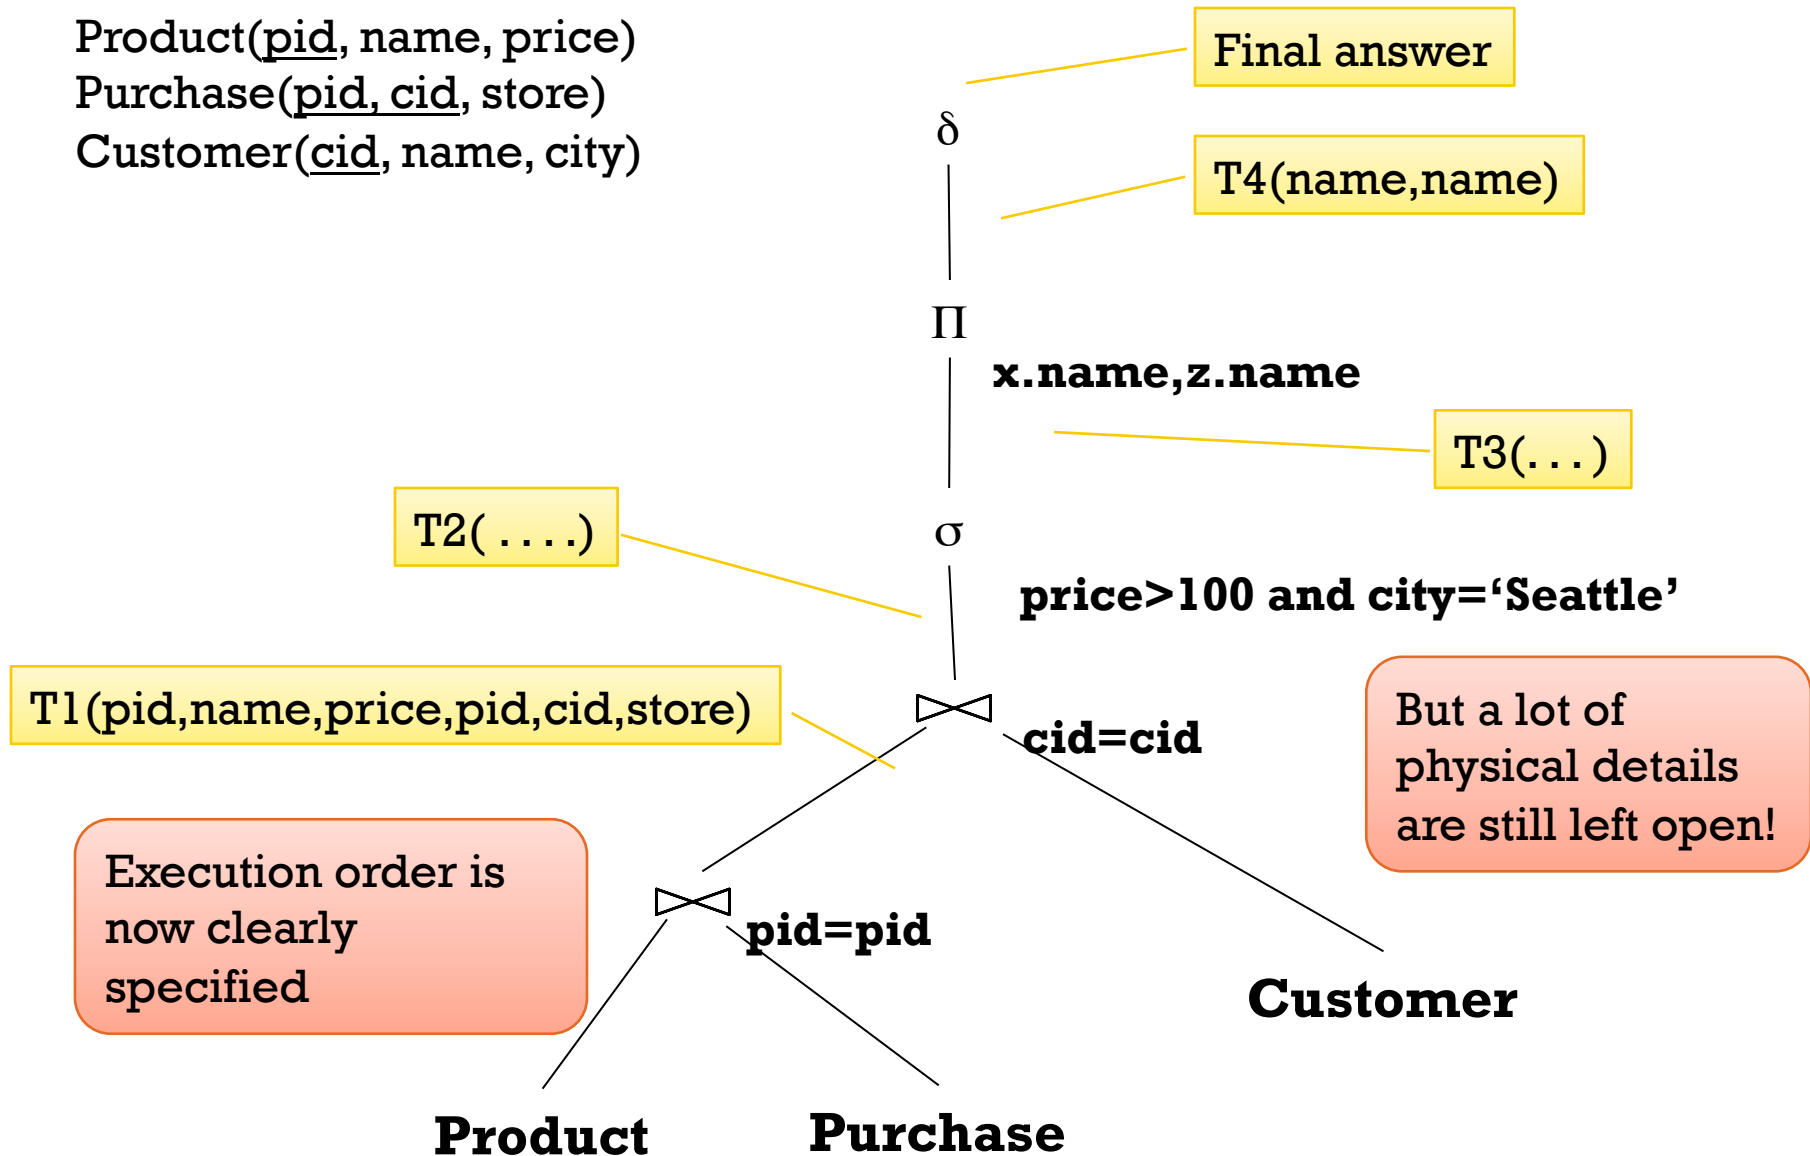
It's clear WHAT we want, unclear HOW to get it

# RELATIONAL ALGEBRA

Product(pid, name, price)

Purchase(pid, cid, store)

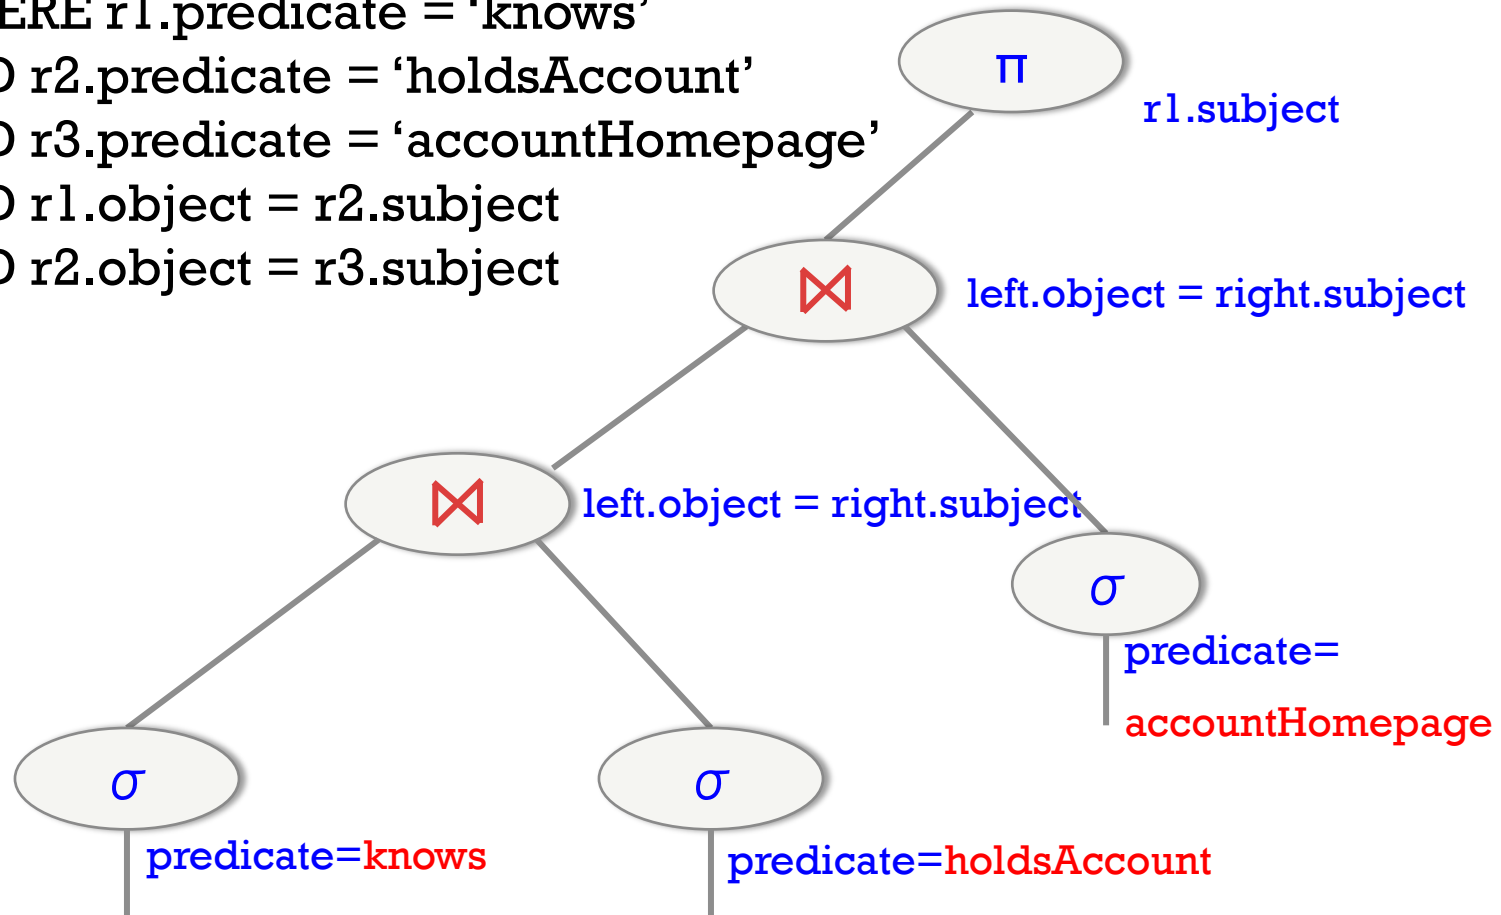
Customer(cid, name, city)




# ANOTHER EXAMPLE

```
SELECT r1.subject
FROM R r1, R r2, R r3
WHERE r1.predicate = 'knows'
AND r2.predicate = 'holdsAccount'
AND r3.predicate = 'accountHomepage'
AND r1.object = r2.subject
AND r2.object = r3.subject
```

R(subject, predicate,  
object)




# KEY IDEA: “LOGICAL DATA INDEPENDENCE”



views

```
SELECT *  
FROM my_sequences
```


*logical data independence*



relations

```
SELECT seq  
FROM ncbi_sequences  
WHERE seq =  
'GATTACGATATTA' ;
```

*physical data independence*



files and  
pointers

```
f = fopen( 'table_file' );  
fseek(10030440);  
while (True) {  
    fread(&buf, 1, 8192, f);  
    if (buf == GATTACGATATTA) {  
        . . .
```

# WHAT ARE VIEWS?

**A *view* is just a query with a name**

**We can use the view just like a real table**

Why can we do  
this?

Because we know that every  
query returns a relation:  
We say that the language is  
“algebraically closed”

# VIEW EXAMPLE

**A view is a relation defined by a query**

Purchase(customer, product, store)    StorePrice(store, price)  
Product(pname, price)

```
CREATE VIEW StorePrice AS  
SELECT x.store, y.price  
FROM Purchase x, Product y  
WHERE x.pid = y.pid
```

This is like a new table  
StorePrice(store,price)

# HOW TO USE A VIEW?

Customer(cid, name, city)

Purchase(customer, product, store)

Product(pname, price)

**A "high end" store is a store that sold some product over 1000. For each customer, find all the high end stores that they visit. Return a set of (customer-name, high-end-store) pairs.**

```
SELECT DISTINCT z.name, u.store
FROM Customer z, Purchase u, StorePrice v
WHERE z.cid = u.customer
AND u.store = v.store
AND v.price > 1000
```



# KEY IDEA: INDEXES

**Databases are especially, but not exclusively, effective at “Needle in Haystack” problems:**

- **Extracting small results from big datasets**
- **Transparently provide “old style” scalability**
- **Your query will *always*\* finish, regardless of dataset size.**
- **Indexes are easily built and automatically used when appropriate**

```
CREATE INDEX seq_idx ON sequence(seq) ;
```

```
SELECT seq  
  FROM sequence  
 WHERE seq = 'GATTACGATATTA' ;
```

***\*almost***

# IN-DATABASE ANALYTICS

BILL HOWE, PHD

DIRECTOR OF RESEARCH, SCALABLE DATA ANALYTICS

UNIVERSITY OF WASHINGTON ESCIENCE INSTITUTE

“There is no point in bringing data ... into the data warehouse environment without integrating it. If the data arrives at the data warehouse in an unintegrated state, it cannot be used to support a corporate view of data. And a corporate view of data is one of the essences of the architected environment.”

- Inmon, 2005, “Building the Data Warehouse”


*Not a good fit for analytics!*

# MATRIX ADDITION, VECTOR REPRESENTATION

*A(row\_number, row float[])*  
*B(row\_number, row float[])*

SELECT A.row\_number, A.vector + B.vector  
FROM A, B  
WHERE A.row\_number = B.row\_number;

# MATRIX MULTIPLICATION



# MATRIX MULTIPLICATION, VECTOR REPRESENTATION

*A(row\_number, row float[])*

SELECT 1, array\_accum(row\_number, vector\*v) FROM A;

$$\vec{x} \cdot \vec{y} = \sum_i x_i y_i$$

# MATRIX TRANSPOSE, VECTOR REPRESENTATION

*A(row\_number, row float[])*

```
SELECT S.col_number, array_accum(A.row_number, A.vector[S.col_number])  
FROM A, generate_series(1,3) AS S(col_number)  
GROUP BY S.col_number;
```

`generate_series` is a *table function*

# MATRIX MULTIPLICATION, SPARSE REPRESENTATION

*A(row\_number, column\_number, value)*  
*B(row\_number, column\_number, value)*

```
SELECT A.row_number, B.column_number, SUM(A.value * B.value)
FROM A, B
WHERE A.column_number = B.row_number
GROUP BY A.row_number, B.column_number
```



# ASIDE: USER-DEFINED FUNCTIONS AND TYPES

## Scalar functions

```
SELECT myfunc(r.a, r.b) ...
```

## Aggregate functions

```
SELECT concat(r.s) ...
```

## Table functions

```
SELECT ... FROM tablefunc(a,b)
```

# EXPERIMENT DESIGN

```
CREATE VIEW design AS
SELECT a.trial_id,
floor (100 * random()) AS row_id
FROM generate_series(1,10000) AS a (trial_id),
generate_series(1,3) AS b (subsample_id)
```

```
CREATE VIEW trials AS
SELECT d.trial_id, AVG(a.values) AS avg_value
FROM design d, T
WHERE d.row_id = T.row_id
GROUP BY d.trial_id
```

Prior to the implementation of this functionality within the DBMS, one Greenplum customer was accustomed to calculating the OLS by exporting data and importing the data into R for calculation, a process that took several hours to complete. They reported significant performance improvement when they moved to running the regression within the DBMS. Most of the benefit derived from running the analysis in parallel close to the data with minimal data movement.

SLIDES CAN BE FOUND AT:  
[TEACHINGDATASCIENCE.ORG](https://teachingdatascience.org)

